

Back-Propagation

Lecture 0 2

Yann Le Cun

Facebook AI Research,

Center for Data Science, NYU


Courant Institute of Mathematical Sciences, NYU

<http://yann.lecun.com>




State Variables

Y LeCun
MA Ranzato



the internal state of the network will be kept in a “state” class that contains two scalars, vectors, or matrices: (1) the state proper, (2) the derivative of the energy with respect to that state.

The input vector is multiplied by the weight matrix.



- fprop: $X_{out} = W X_{in}$

- bprop to input:

$$\frac{\partial E}{\partial X_{in}} = \frac{\partial E}{\partial X_{out}} \frac{\partial X_{out}}{\partial X_{in}} = \frac{\partial E}{\partial X_{out}} W$$

- by transposing, we get column vectors:

$$\frac{\partial E}{\partial X_{in}}' = W' \frac{\partial E}{\partial X_{out}}'$$

- bprop to weights:


$$\frac{\partial E}{\partial W_{ij}} = \frac{\partial E}{\partial X_{out i}} \frac{\partial X_{out i}}{\partial W_{ij}} = X_{in j} \frac{\partial E}{\partial X_{out i}}$$

- We can write this as an outer-product:

$$\frac{\partial E}{\partial W}' = \frac{\partial E}{\partial X_{out}}' X_{in}'$$

Tanh module (or any other pointwise function)


Y LeCun
MA Ranzato



- fprop: $(X_{out})_i = \tanh((X_{in})_i + B_i)$
- bprop to input:
$$\left(\frac{\partial E}{\partial X_{in}}\right)_i = \left(\frac{\partial E}{\partial X_{out}}\right)_i \tanh'((X_{in})_i + B_i)$$
- bprop to bias:
$$\frac{\partial E}{\partial B_i} = \left(\frac{\partial E}{\partial X_{out}}\right)_i \tanh'((X_{in})_i + B_i)$$
- $$\tanh(x) = \frac{2}{1 + \exp(-x)} - 1 = \frac{1 - \exp(-x)}{1 + \exp(-x)}$$

Euclidean Distance Module (Squared Error)


Y LeCun
MA Ranzato



- fprop: $X_{out} = \frac{1}{2} \|X_{in} - Y\|^2$
- bprop to X input: $\frac{\partial E}{\partial X_{in}} = X_{in} - Y$
- bprop to Y input: $\frac{\partial E}{\partial Y} = Y - X_{in}$

Y connector and Addition modules

Y LeCun
MA Ranzato



- The PLUS module: a module with K inputs X_1, \dots, X_K (of any type) that computes the sum of its inputs:


$$X_{\text{out}} = \sum_k X_k$$


$$\text{back-prop: } \frac{\partial E}{\partial X_k} = \frac{\partial E}{\partial X_{\text{out}}} \quad \forall k$$

- The BRANCH module: a module with one input and K outputs X_1, \dots, X_K (of any type) that simply copies its input on its outputs:

$$X_k = X_{\text{in}} \quad \forall k \in [1..K]$$

$$\text{back-prop: } \frac{\partial E}{\partial \text{in}} = \sum_k \frac{\partial E}{\partial X_k}$$





- A module with K inputs X_1, \dots, X_K (of any type) and one additional discrete-valued input Y .
- The value of the discrete input determines which of the N inputs is copied to the output.

$$X_{\text{out}} = \sum_k \delta(Y - k) X_k$$

$$\frac{\partial E}{\partial X_k} = \delta(Y - k) \frac{\partial E}{\partial X_{\text{out}}}$$


the gradient with respect to the output is copied to the gradient with respect to the switched-in input. The gradients of all other inputs are zero.

SoftMax Module (should really be called SoftArgMax)

Y LeCun
MA Ranzato

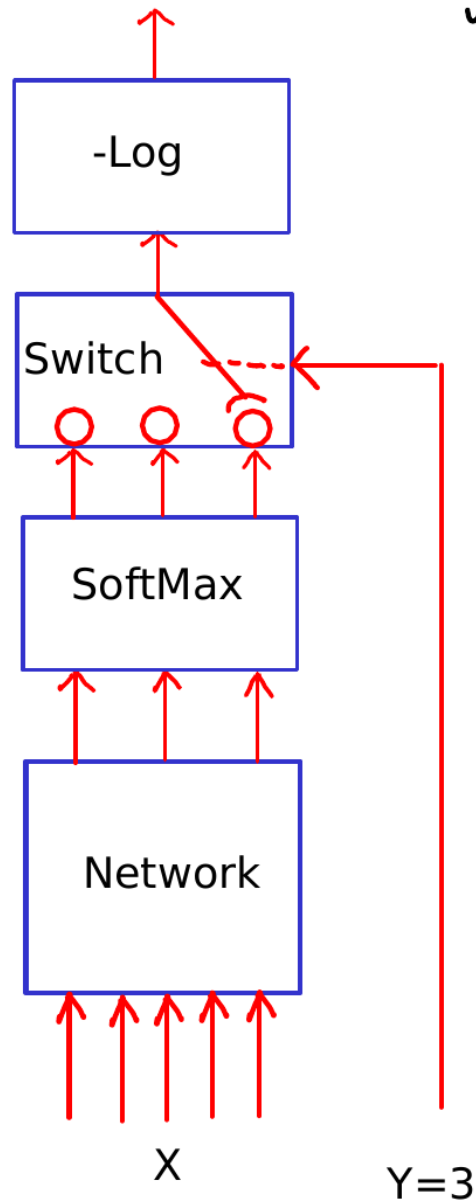
- Transforms scores into a discrete probability distribution
 - Positive numbers that sum to one.
- Used in multi-class classification

$$p_k = \frac{e^{\beta x_k}}{\sum_j e^{\beta x_j}}$$




SoftMax Module: Loss Function for Classification

Y LeCun
MA Ranzato




• $-\text{LogSoftMax}: -\frac{1}{\beta} \log p_k = -X_k + \frac{1}{\beta} \log \sum_j e^{\beta X_j}$

- Maximum conditional likelihood
- Minimize $-\log$ of the probability of the correct class.

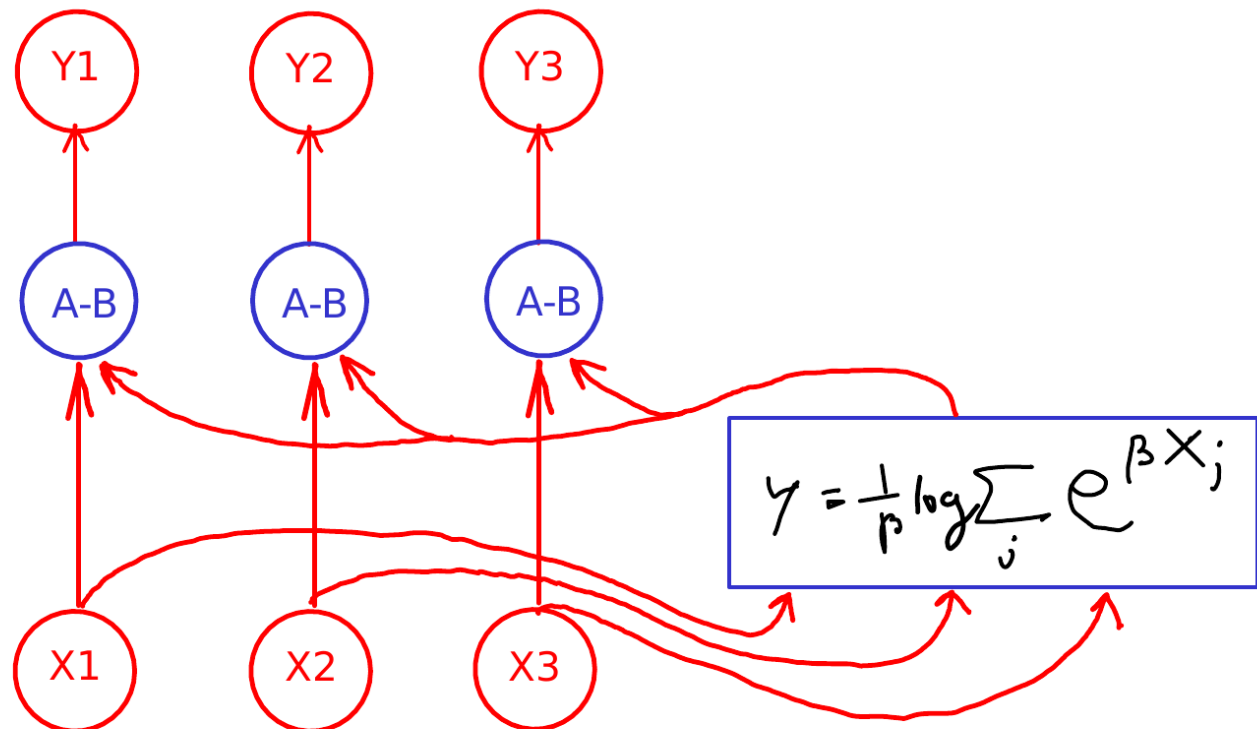


LogSoftMax Module

Y LeCun
MA Ranzato



- Transforms scores into a discrete probability distribution
- $\text{LogSoftMax} = \text{Identity} - \text{LogSumExp}$



- Log of normalization term for SoftMax

- Fprop

$$X_{out} = \frac{1}{\beta} \sum_j e^{\beta X_j}$$

- Bprop


$$\frac{\partial E}{\partial X_k} = \frac{\partial E}{\partial X_{out}} \cdot \frac{e^{\beta X_k}}{\sum_j e^{\beta X_j}}$$

- Or:

$$\frac{\partial E}{\partial X_k} = \frac{\partial E}{\partial X_{out}} \cdot P_k \quad P_k = \frac{e^{\beta X_k}}{\sum_j e^{\beta X_j}}$$

Backprop works through any modular architecture

Y LeCun
MA Ranzato



Any connection is permissible

- ▶ Networks with loops must be “unfolded in time”.

Any module is permissible

- ▶ As long as it is continuous and differentiable almost everywhere with respect to the parameters, and with respect to non-terminal inputs.

Backprop in Practice

Y LeCun
MA Ranzato

- **Use ReLU non-linearities (tanh and logistic are falling out of favor)**
- **Use cross-entropy loss for classification**
- **Use Stochastic Gradient Descent on minibatches**
- **Shuffle the training samples**
- **Normalize the input variables (zero mean, unit variance)**
- **Schedule to decrease the learning rate**
- **Use a bit of L1 or L2 regularization on the weights (or a combination)**
 - ▶ But it's best to turn it on after a couple of epochs
- **Use “dropout” for regularization**
 - ▶ Hinton et al 2012 <http://arxiv.org/abs/1207.0580>
- **Lots more in [LeCun et al. “Efficient Backprop” 1998]**
- **Lots, lots more in “Neural Networks, Tricks of the Trade” (2012 edition) edited by G. Montavon, G. B. Orr, and K-R Müller (Springer)**


Deep Learning is Non-Convex

Y LeCun
MA Ranzato

Example: what is the loss function for the simplest 2-layer neural net ever

- Function: 1-1-1 neural net. Map 0.5 to 0.5 and -0.5 to -0.5 (identity function) with quadratic cost:

$$y = \tanh(W_1 \tanh(W_0 x)) \quad L = (0.5 - \tanh(W_1 \tanh(W_0 0.5))^2$$



The Objective Function of Multi-layer Nets is Non Convex

Y LeCun



MA Ranzato

1-1-1 network

$$Y = W1 * W2 * X$$

Objective: identity function with quadratic

One sample: $X=1, Y=1$ $L(W) = (1 - W1 * W2)$



Solution

Saddle point

Solution

